

Introduction to HDF5
and
a few advanced HDF5 features
for GLCS

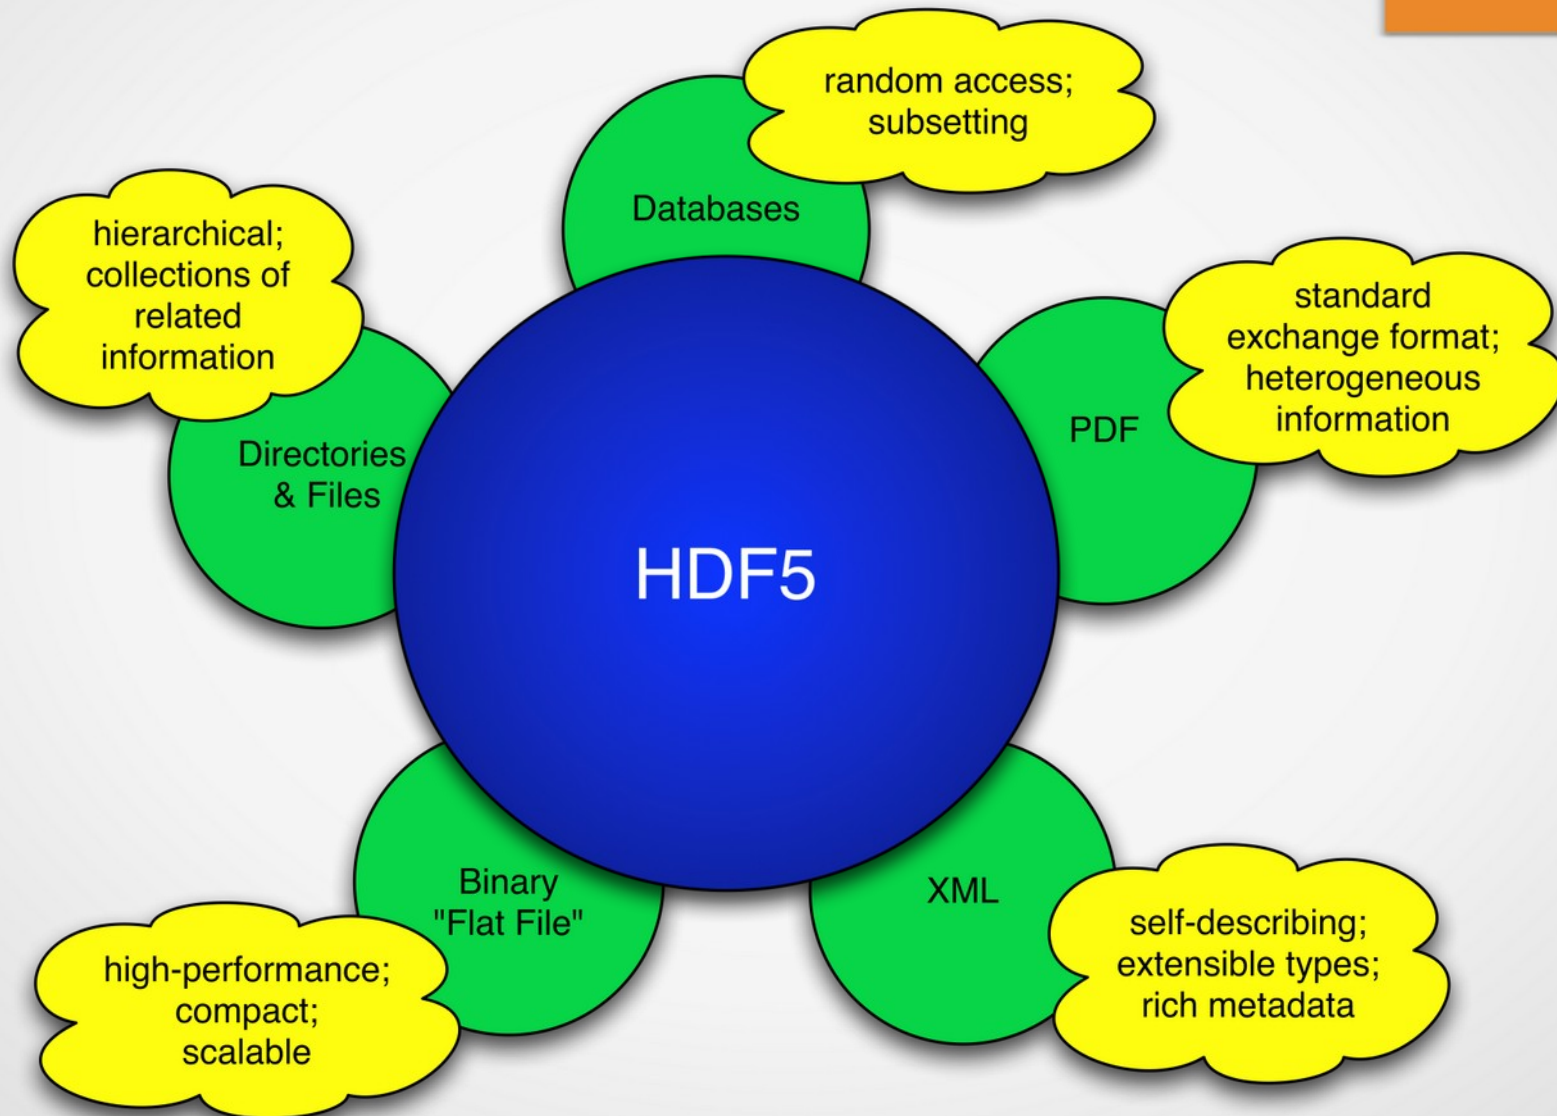


What is HDF5?

- HDF5 == Hierarchical Data Format, v5
- Open file format
 - Designed for high volume or complex data
- Open source software
 - Works with data in the format
- An extensible data model
 - Structures for data organization and specification



HDF5 is like ...





HDF5 Data Model

- An HDF5 file is a container that holds data objects.





HDF5 Dataset


HDF5 Datatype

Integer: 32-bit, LE

HDF5 Dataspace

Rank	Dimensions
3	Dim[0] = 4
	Dim[1] = 5
	Dim[2] = 7

Specifications for single data element and array dimensions



Multi-dimensional array of identically typed data elements

- HDF5 datasets organize and contain data elements.
- HDF5 datatype describes individual data elements.
- HDF5 dataspace describes the logical layout of the data elements.




HDF5 Dataspace

- Describes the logical layout of the elements in an HDF5 dataset
 - NULL
 - no elements
 - Scalar
 - single element
 - Simple array (most common)
 - multiple elements organized in a rectangular array
 - rank = number of dimensions
 - dimension sizes = number of elements in each dimension
 - maximum number of elements in each dimension
 - may be fixed or unlimited

HDF5 Dataspace

Two roles:

- Dataspace contains spatial information
 - Rank and dimensions
 - Permanent part of dataset definition



Rank = 2

Dimensions = 4x6

- Partial I/O: Dataspace describes application's data buffer and data elements participating in I/O



Rank = 1

Dimension = 10




HDF5 Datatypes

- Describe individual data elements in an HDF5 dataset
- Wide range of datatypes supported
 - Integer
 - Float
 - Enum
 - Array
 - User-defined (e.g., 13-bit integer)
 - Variable-length types (e.g., strings, vectors)
 - Compound (similar to C structs)
 - More ...



HDF5 Dataset



- Datatype: 32-bit Integer
- Dataspace:
 - Rank = 2
 - Dimensions = 5 x 3



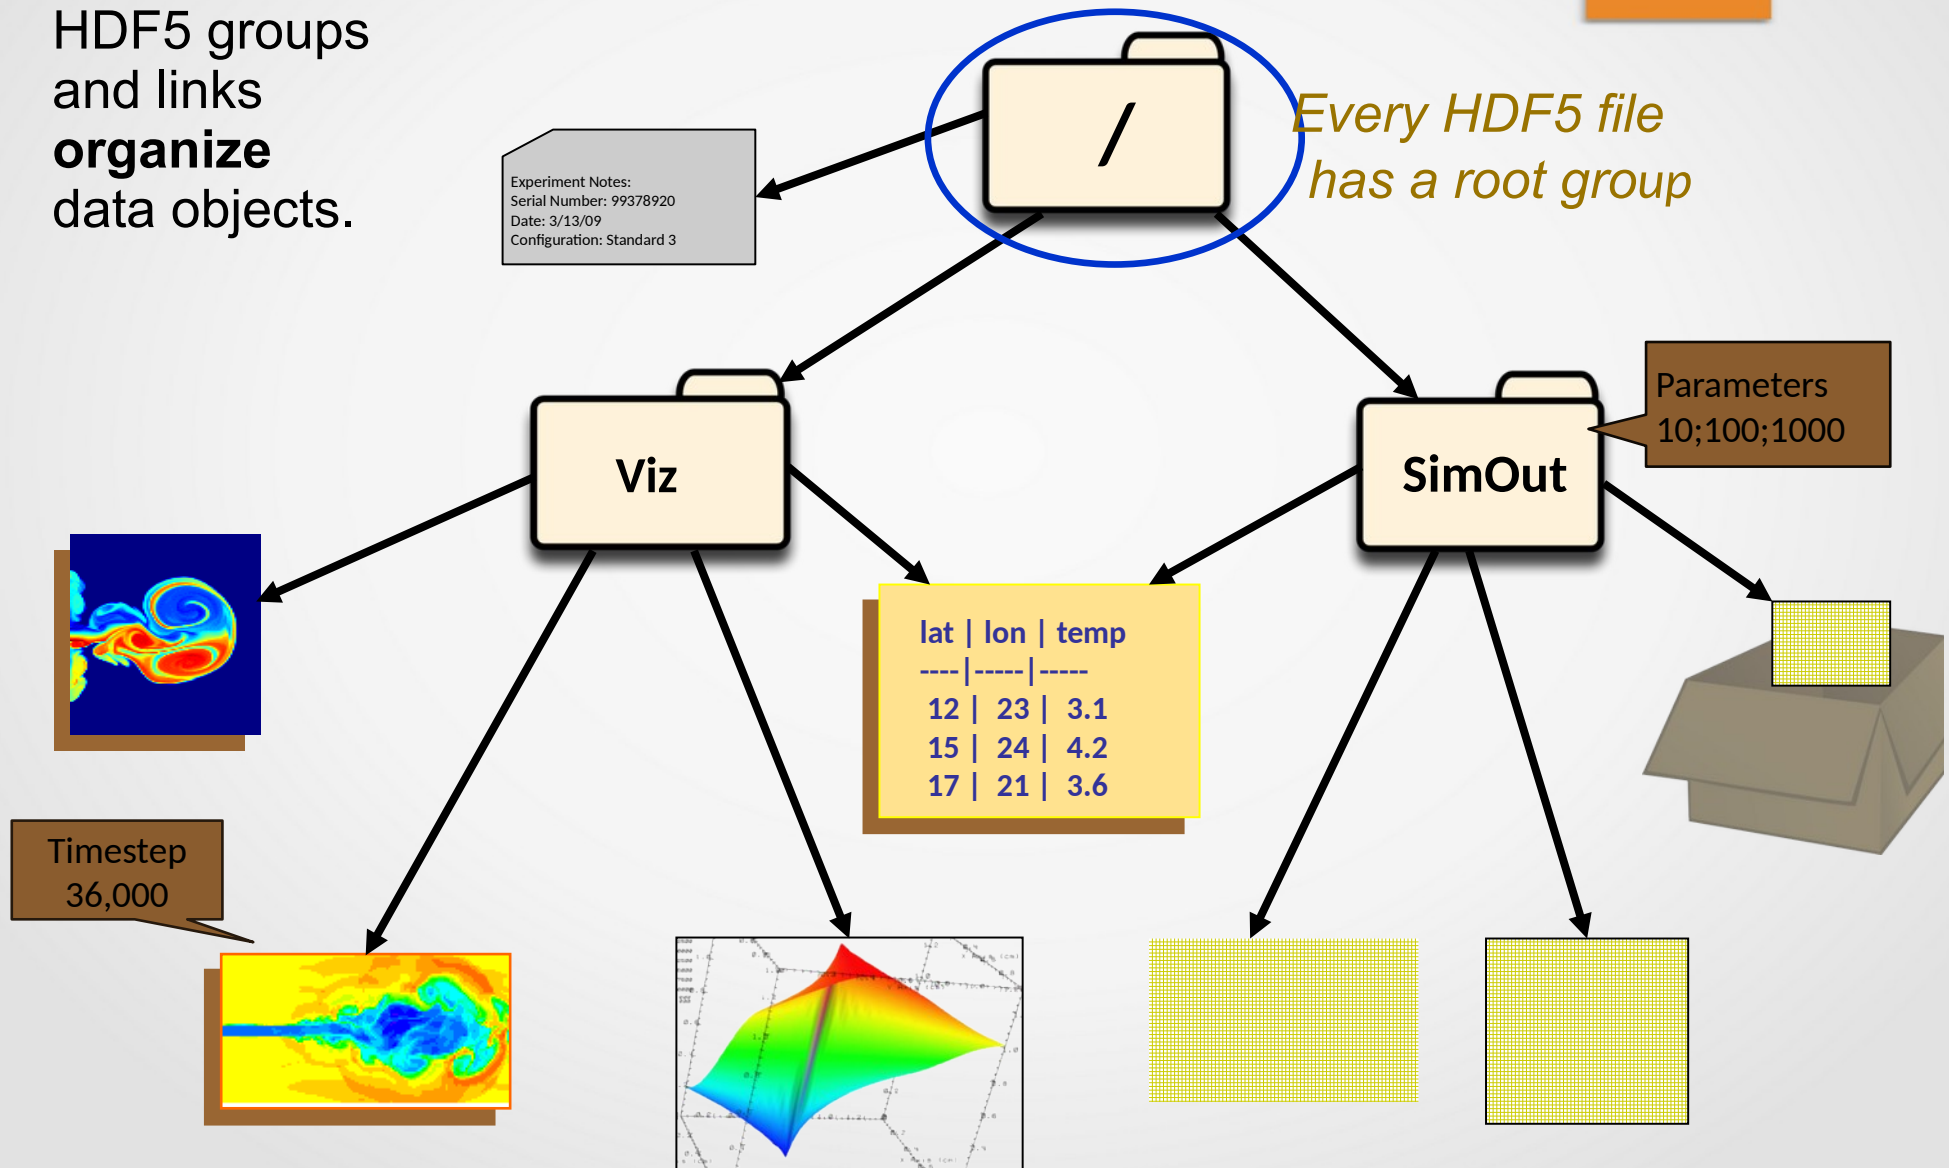
HDF5 Attributes

- Typically contain user metadata
- Have a name and a value
- Attributes “decorate” HDF5 objects
- Value is described by a datatype and a dataspace
- Analogous to a dataset, but
 - do not support partial I/O operations
 - can not be compressed or extended



HDF5 Groups and Links

HDF5 groups and links **organize** data objects.





Useful Tools

- h5dump:
 - Tool to “dump” or display contents of HDF5 files
- h5cc, h5c++, h5fc:
 - Scripts to compile applications
- HDFView:
 - Java browser to view HDF5 files
 - <http://www.hdfgroup.org/hdf-java-html/hdfview/>



The General HDF5 API

➤ Prefix: H5?

- ? is a character corresponding to the type of object the function acts on

➤ Example:

- H5D: Dataset interface *e.g.*, H5Dread
- H5F: File interface *e.g.*, H5Fopen
- H5S: DataSpace interface *e.g.*, H5Sclose



The HDF5 API

- For flexibility, the API is extensive
 - 300+ functions
- This can be daunting... but there is hope
 - A few functions can do a lot
 - Start simple
 - Build up knowledge as more features are needed



General Programming Paradigm

- Object is opened or created
 - Object is accessed, possibly many times
 - Object is closed
-
- Properties of object are optionally defined
 - Creation properties (e.g., use chunking storage)
 - Access properties



Basic Functions

H5**F**create (H5**F**open)

create (open) File

H5**S**create_simple/H5**S**create *create dataSpace*

H5**D**create (H5**D**open)

create (open) Dataset

H5**D**read, H5**D**write

access Dataset

H5**D**close

close Dataset

H5**S**close

close dataSpace

H5**F**close

close File



Other Common Functions

➤ Data**S**paces:

- H5**S**elect_hyperslab (Partial I/O)
- H5**S**elect_elements (Partial I/O)
- H5Dget_**s**pace

➤ **G**roups:

- H5**G**create, H5**G**open, H5**G**close

➤ **P**roperty lists:

- H5**P**create, H5**P**close



Code: Create a File


```
hid_t file_id = H5Fcreate("file.h5", H5F_ACC_TRUNC,  
                           H5P_DEFAULT,H5P_DEFAULT);  
// ...  
H5Fclose (file_id);
```

“/” (root)



Code: Create a Dataset

```
hid_t file_id = H5Fcreate("file.h5", H5F_ACC_TRUNC,  
                           H5P_DEFAULT, H5P_DEFAULT);  
  
hsize_t dims[2] = {4, 6};  
hid_t dataspace_id = H5Screate_simple(2, dims, NULL);  
  
hid_t dataset_id = H5Dcreate(file_id, "A", H5T_STD_I32BE,  
                             dataspace_id, H5P_DEFAULT,  
                             H5P_DEFAULT, H5P_DEFAULT);  
  
// ...  
  
H5Dclose(dataset_id);  
  
H5Sclose(dataspace_id);  
  
H5Fclose(file_id);
```





Code: Create a Group


```
// ...
```

```
hid_t file_id = H5Fopen("file.h5", H5F_ACC_RDWR,  
                        H5P_DEFAULT);
```

```
hid_t group_id = H5Gcreate(file_id, "B", H5P_DEFAULT,  
                           H5P_DEFAULT, H5P_DEFAULT);
```

```
// ...
```

```
H5Gclose(group_id);  
H5Fclose(file_id);
```





Output of h5dump

```
$ h5dump file.h5
```

```
HDF5 "file.h5" {  
  GROUP "/" {  
    DATASET "A" {  
      DATATYPE  H5T_STD_I32BE  
      DATASPACE  SIMPLE { ( 4, 6 ) / ( 4, 6 ) }  
      DATA {  
        (0,0): 0, 0, 0, 0, 0, 0,  
        (1,0): 0, 0, 0, 0, 0, 0,  
        (2,0): 0, 0, 0, 0, 0, 0,  
        (3,0): 0, 0, 0, 0, 0, 0  
      }  
    }  
    GROUP "B" {  
  }  
}
```



Example Code - H5Dwrite

```
int wdata[4][6];  
for (i = 0; i < 4; i++)  
    for (j = 0; j < 6; j++)  
        wdata[i][j] = i * 6 + j + 1;
```

```
// ...
```

```
H5Dwrite(dataset_id, H5T_NATIVE_INT, H5S_ALL, H5S_ALL,  
         H5P_DEFAULT, wdata);
```



Output of h5dump after writing

```
$ h5dump file.h5
```

```
HDF5 "file.h5" {  
  GROUP "/" {  
    DATASET "A" {  
      DATATYPE  H5T_STD_I32BE  
      DATASPACE  SIMPLE { ( 4, 6 ) / ( 4, 6 ) }  
      DATA {  
        (0,0): 1, 2, 3, 4, 5, 6,  
        (1,0): 7, 8, 9, 10, 11, 12,  
        (2,0): 13, 14, 15, 16, 17, 18,  
        (3,0): 19, 20, 21, 22, 23, 24  
      }  
    }  
    GROUP "B" {  
    }  
  }  
}
```



How to write a row?

```
$ h5dump file.h5
```

```
HDF5 "file.h5" {  
  GROUP "/" {  
    DATASET "A" {  
      DATATYPE  H5T_STD_I32BE  
      DATASPACE  SIMPLE { ( 4, 6 ) / ( 4, 6 ) }  
      DATA {  
        (0,0): 0, 0, 0, 0, 0, 0,  
        (1,0): 7, 8, 9, 10, 11, 12,  
        (2,0): 0, 0, 0, 0, 0, 0,  
        (3,0): 0, 0, 0, 0, 0, 0  
      }  
    }  
    GROUP "B" {  
  }  
}
```




Describe a Subset in HDF5

- Before writing and reading a subset of data one has to describe it to the HDF5 Library
- HDF5 APIs and documentation refer to a subset as a “selection” or “hyperslab selection”
- If specified, HDF5 Library will perform I/O on a selection only and not on all elements of a dataset.




Types of Selections in HDF5

- Two types of selections
 - Hyperslab selection
 - Regular hyperslab
 - Simple hyperslab
 - Result of set operations on hyperslabs (union, difference, ...)
 - Point selection
- Hyperslab selection is especially important for doing parallel I/O in HDF5




Contiguous subset or sub-array



Collection of regularly spaced blocks of equal size

Hyperslab Selection




Result of union operation on three simple hyperslabs



HDF5 Hyperslab Description

- Everything is “measured” in number of elements
 - Uses row-major ordering (C order) for coordinates
 - Example:
 - Start - starting location of a hyperslab (1,1)
 - Stride - number of elements that separate each block (3,2)
 - Count - number of blocks (2,6)
 - Block - block size (2,1)






Simple Hyperslab Description

Two ways to describe a simple hyperslab


➤ As several blocks

- Stride – (1,1)
- Count – (2,6)
- Block – (2,1)



➤ As one block

- Stride – (1,1)
- Count – (1,1)
- Block – (4,6)






Writing a row

- Memory space selection is 1-dim array of size 6



- File space selection

- $\text{start} = \{1,0\}$, $\text{stride} = \{1,1\}$, $\text{count} = \{1,6\}$, $\text{block} = \{1,1\}$



- Number of elements selected in memory must be the same as selected in the file



Writing a row

```
// ...  
hsize_t dims[1] = {6};  
hid_t mspace_id = H5Screate_simple(1, dims, NULL);  
  
hid_t fspace_id = H5Dget_space(dataset_id);  
  
hsize_t start[2] = {1, 0};  
hsize_t count[2] = {1, 6};  
H5Sselect_hyperslab(fspace_id, H5S_SELECT_SET, start, NULL, count,  
                    NULL);  
  
H5Dwrite(dataset_id, H5T_NATIVE_INT, mspace_id, fspace_id,  
         H5P_DEFAULT, wdata);
```

Parallel HDF5

- Open files with a MPI communicator
 - Most file operations become a MPI collective
 - All parts of the file are accessible by all processes
 - All objects in the file are accessible by all processes
 - Multiple processes can write to the same dataset
- Use partial I/O to write your local data block
 - But no overlap !



//HDF5 : Opening a file

For all processes of the MPI communicator :

- Set up an access template object
 - `hid_t H5Pcreate(H5P_FILE_ACCESS)`
- Set up the access template for parallel I/O
 - `H5Pset_fapl_mpio(hid_t fapl_id,
MPI_Comm comm,
MPI_Info info)`
 - `comm` & `info` passed to `MPI_FILE_OPEN`
 - `Info` : usually `MPI_INFO_NULL`
- Open the file using the access template
- Close the file



Parallel HDF5 : Example

```
hid_t plist_id = H5Pcreate(H5P_FILE_ACCESS);  
                H5Pset_fapl_mpio(plist_id, MPI_COMM_WORLD,  
                                MPI_INFO_NULL);  
  
hid_t file_id = H5Fcreate(H5FILE_NAME, H5F_ACC_TRUNC, H5P_DEFAULT,  
                          plist_id);  
  
// ...  
  
H5Pclose(plist_id);  
  
H5Fclose(file_id);  
  
MPI_Finalize();
```



Parallel dataset access

- Create a file transfer property list
 - `hid_t H5Pcreate(H5P_DATASET_XFER)`
- Set the data transfer mode to independent I/O collective I/O
 - `H5Pset_dxpl_mpio (hid_t dxpl_id,
H5FD_mpio_xfer_t xfer_mode)`
 - `dxpl_id`: Data transfer property list identifier
 - `xfer_mode`: Transfer mode:
 - `H5FD_MPIO_INDEPENDENT`: independent I/O access
 - `H5FD_MPIO_COLLECTIVE`: collective I/O access
- Access the dataset with the defined transfer property list
 - All processes that have opened a dataset may do collective I/O.
 - Each process may do an independent and arbitrary number of data I/O access calls
 - `H5Dwrite` / `H5Dread`



Parallel access: example

```
// ...  
hid_t filespace = H5Screate_simple(RANK, dims, NULL);  
hid_t dset_id = H5Dcreate(file_id, DATASETNAME, H5T_NATIVE_INT,  
                           filespace, H5P_DEFAULT, H5P_DEFAULT,  
                           H5P_DEFAULT);  
hid_t plist_id = H5Pcreate(H5P_DATASET_XFER);  
H5Pset_dxpl_mpio(plist_id, H5FD_MPIO_COLLECTIVE);  
  
H5Dwrite(dset_id, H5T_NATIVE_INT, H5S_ALL, H5S_ALL, plist_id, data);  
  
H5Dclose(dset_id);  
H5Sclose(filespace);  
H5Pclose(plist_id);
```



Your turn now

- Get the TD description :

https://raw.githubusercontent.com/jbigot/glcs_2020-2021/master/td1.pdf

- Set-up your environment :

- Install docker

```
sudo apt install docker.io  
sudo gpaswd -a $USER docker
```

- Get the TD environment script

```
wget https://raw.githubusercontent.com/jbigot/glcs_2020-  
2021/master/glcs-td1.sh
```

- Generate the TD directory

```
bash glcs-td1.sh
```